

Inlining in the Jive Compiler Backend

Christian Chavez

Thursday 7th May, 2015

Nico Reissmann, Magnus Jahre, and Christian Chavez are with the Norwegian University of Science and Technology (NTNU).

Abstract

Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Contents

1	Introduction	4
2	The Regionalized Value State Dependence Graph	6
2.1	Edges	6
2.2	Nodes	7
2.2.1	Simple Nodes	7
2.2.2	Complex Nodes	8
3	The Inliner	13
3.1	Deciding which recursive functions to inline	13
3.2	The order of call sites inlined	14
3.3	Inlining a call site	15
4	Methodology	16
4.1	The Inlining Conditions (ICs)	16
5	Results	17
6	Related Work	18
7	Conclusion	20
8	Further Work	21
8.1	Dynamic profiling and adaptive compilation	21
8.2	Choosing loop breakers more carefully	21
8.3	The ordering of which call sites to order first	21
8.3.1	Something	21
	List of Figures	22
	List of Listings	23
9	References	24
	Appendices	25
A	Project Description	25

Todo's present in document:

Find reference, if time allows. Nico says to look in Dragon book, figure out where you can borrow that one.	4
There seems to be a big paper named "Variable Subsumption with Constant Folding" by a "Ken, Kennedy", but I can't find the paper. It's from 1973-1974. Checked acm.org, citeseerx, google scholar. Should I cite it without having read it? I would think not...	4
Finish describing layout/outline of paper. What do the remaining sections discuss in turn?	5
Where in the report does Nico want this?	5
Make a figure or three of ϕ -regions, one with a self-recursive lambda, one with some mutually recursive ones (the example from PVV w/Torje), and perhaps also one based on GHC papers FGHPQ example.	13
Because of above paragraph, which I think belongs to Section 3.3, ask Nico if he insists that Section 3.2 should come before Section 3.3. I feel that this makes little sense since the one that is first now relies on the second...	14
Describe the algorithm and inliner conditions we land for evaluating a call site on after testing.	15
Need an introduction here, no?	16
Make a subsection or paragraph showing the results of <i>apply</i> -nodes which link to static calls, vs all calls.	17
Basically re-iterate Waterman's PhD idea.	21
Bas' MSc. idea	21
Lots to say on ordering of call sites to be inlined	21
If time permits, I want to briefly discuss the idea of evaluating all the apply nodes (in order), before inlining any. So as to be able to see the total cost of inlining a successive chain of function calls in one direction or the other.	21

1 Introduction

Since the 1950s, compilers have been translating higher-level programming languages into machine languages. The purpose of a compiler is two-fold: translate human-readable code into machine language, and optimize the translated code. Compilers use many code optimization techniques, such as *Common Subexpression Elimination* (CSE) and *Dead Code Elimination* (DCE).

Find reference, if time allows. Nico says to look in Dragon book, figure out where you can borrow that one.

There seems to be a big paper named “Variable Subsumption with Constant Folding” by a “Ken, Kennedy”, but I can’t find the paper. It’s from 1973-1974. Checked acm.org, citeseerx, google scholar. Should I cite it without having read it? I would think not...

Another code optimization is *inlining*, which replaces the call site of a function with its body. Listing 1 shows the definition of functions A() and B(). If A() is inlined into B, the body of B() becomes `return y + 3 + 2;`, allowing *Constant Folding* (CF) to replace `3+2` with `5`.

```
int A(int x){
    return x + 3;
}

int B(int y){
    return A(y) + 2;
}
```

Listing 1: C/C++ code showing the definitions of A() and B() when exemplifying CF.

The benefits of inlining are mainly two-fold: The first one is the removal of function call overhead. Function call overhead is the cost in memory needed on the stack, as well as the CPU cycles needed for setting up and performing the call. The second is the potential for unveiling the application of additional optimizations as demonstrated in Listing 1.

The drawbacks of inlining are code-duplication, and an increased compile time. In specific situations, work-duplication can also occur [10]. Listing 2 exemplifies a situation where inlining can lead to code-duplication, if C() is inlined into D(). Code-duplication can occur since C() is called more than once in D(). This is due to the big expression `e` in C() is copied into D() as soon as one or more of invocations of C() are inlined. However, code-duplication might be mitigated if CSE is applicable. Also, if C() is static with all its invocations inlined, then the definition of C() may be removed through the application of DCE.

```
static int C(int a){
    return e; //Big expression, depending on a
}

int D(int x, int y){
    return C(x) + C(y);
}
```

Listing 2: C/C++ code showing the definitions of C() and D(), when exemplifying code-duplication.

If inlining is performed blindly on all function call sites, non-termination of the compilation can occur. This can happen when the compiler attempts to inline recursive functions. Hence, recursive functions need to be handled carefully. The literature proposes mainly two approaches for inlining recursive functions:

1. Avoid non-termination of the compilation by only inlining recursive functions to a certain depth [10][11], and therefore breaking the recursive cycle.
2. In a mutually recursive environment, decide on one or more functions to be *loop breakers*, and mark them to never be inlined. Loop breakers are chosen so that the recursive call cycle will be broken in the mutually recursive environment. Having chosen correct loop breakers permits inlining the remaining recursive functions in the mutually recursive environment, *without* risking non-termination of the compiler [7][10].

This report describes the construction of an inliner for the Jive compiler backend, detailing the design and architecture. Jive uses an *intermediate representation* (IR) called the *Regionalized Value State Dependence Graph* (RVSDG).

The RVSDG [1] described in Section 2, is a *demand-based* and *directed acyclic graph* (DAG) where nodes represent computations, and edges represent the dependencies between these computations.

Section 3 explains how the inliner is able to handle recursive functions, and how the inliner permits the configuration of different heuristics to allow rapid exploration of the parameter space. How the RVSDG affects the design of an inliner, and the algorithms used by the heuristics deciding what to inline, are also detailed in Section 3.

Finish describing layout/outline of paper. What do the remaining sections discuss in turn?

Focus is put on whether the RVSDG simplifies or complicates the implementation of the inliner, as well as the impact of the RVSDG on an inliner, and the process of inlining, compared to commonly used IRs.

In Section 4, the implemented inliner is evaluated before we conclude in Section 7. In the evaluation, we focus on how different heuristics have different consequences, such as code-duplication, and others.

Finally, in Section 8 we discuss ideas for potential further research. A detailed description of the project assignment can be found in Appendix A.

Where in the report does Nico want this?

2 The Regionalized Value State Dependence Graph

The *Regionalized Value State Dependence Graph* [1] (RVSDG) is a *directed acyclic demand-based dependence graph*, consisting of nodes representing computations and edges representing the dependencies between nodes. Each node has inputs and outputs connected through edges. The arity and order of inputs and outputs depend on the operation the node represents.

Figure 1 exemplifies how the C/C++ code on the left can be represented as an RVSDG. The nodes in Figure 1 represent operations in a program, while the edges between the nodes show the dependencies nodes have to each other, thus giving the order of execution.

In all RVSDG examples depicted in this report, the order of inputs in a node goes clockwise. The first input of a node is the one closest to the bottom left corner of the node.




Figure 1: Example of an RVSDG subgraph equivalent to the C/C++ on the left.

2.1 Edges

The RVSDG has data dependence edges and state dependence edges, representing data and state dependencies operations have to each other, respectively. An example of data dependence edges are the operands used in an addition, such as in Figure 1.

State dependence edges are used to preserve the semantics of the program when the program has side-effecting operations. If there are no data dependencies between operations, state dependence edges can give the needed order of execution. Figure 2 illustrates the use of state dependence edges, depicted as dotted edges.

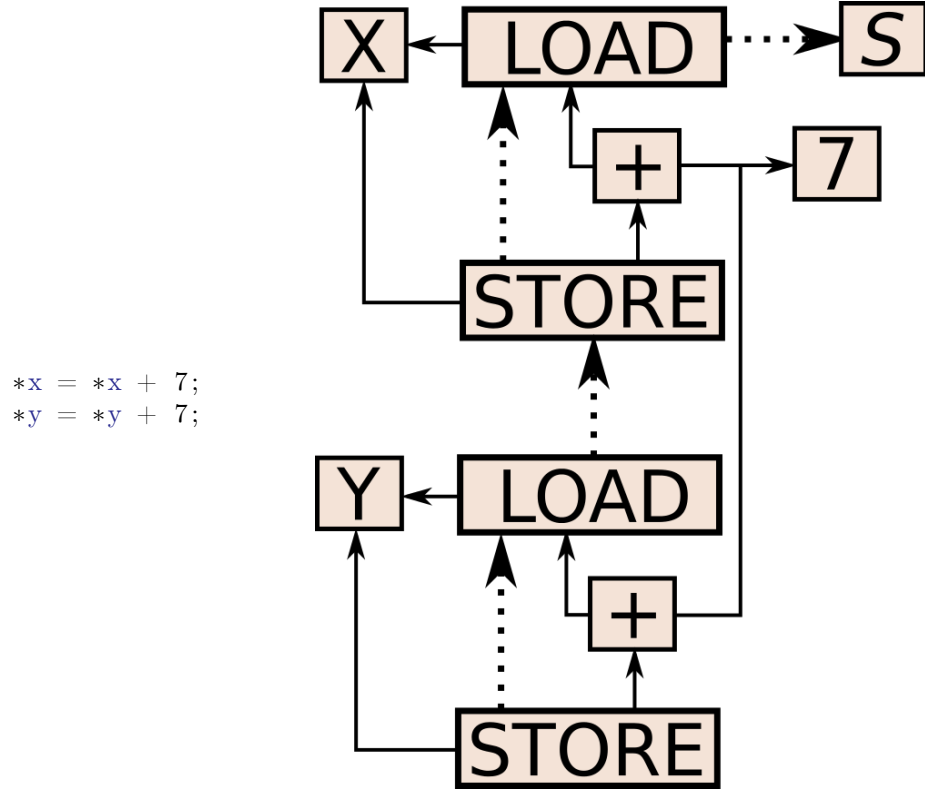


Figure 2: Example of an RVSDG subgraph equivalent to the C/C++ on the left.

The purpose of the *S*-node in Figure 2 is being a placeholder for the state-dependence edge of the load operation of *x*.

2.2 Nodes

The RVSDG has two kinds of nodes: simple nodes and complex nodes. The arity and order of inputs and outputs of any RVSDG node need to match the operation.

2.2.1 Simple Nodes

Simple nodes are used to represent primitive operations, such as addition and subtraction. Figure 1 is an example of an RVSDG containing only simple nodes.

One simple node of special interest for this report is the *apply*-node. An *apply*-node represents the call site of a function. The first input argument of an *apply*-node is the function the *apply*-node invokes. The remaining inputs are the arguments to this function. Likewise, its results are the results of the invocation of its function. Order and arity of inputs and outputs need to match the arguments and results of the function, respectively.

2.2.2 Complex Nodes

Complex nodes contain one or more RVSDG subgraphs, which is why they are also referred to as *regions*. Differing from the simple nodes with their contained subgraph, complex nodes may besides the normal inputs and outputs, also have internal inputs and outputs. Figure 3 shows which inputs/outputs are the external ones, and which are the internal ones. Figure 3 also illustrates how the values of the external inputs are mapped to the internal outputs of each subregion, and vica versa with each subregion's internal inputs being mapped to the external outputs.

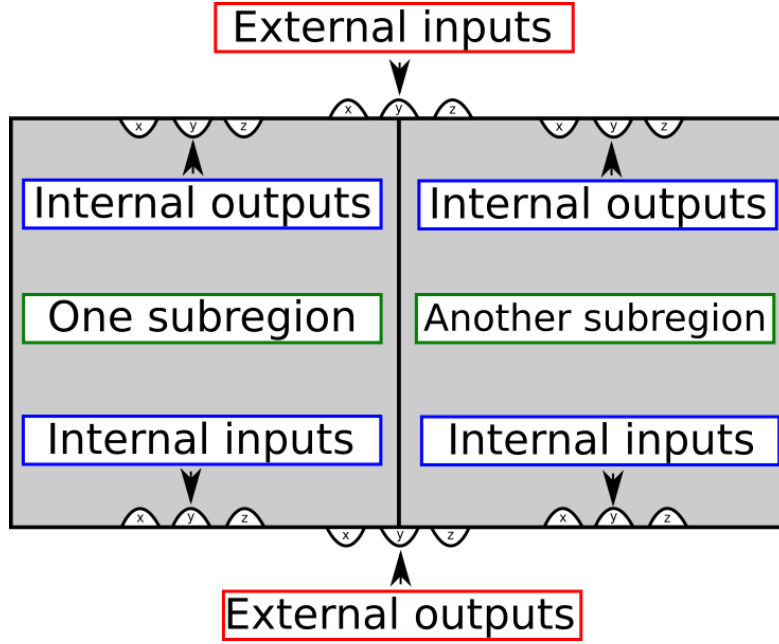


Figure 3: A minimal example of a complex node, showing which inputs/outputs are external/internal, and how they can have multiple subregions.

As Figure 3 also shows, some complex nodes may also have multiple *subregions*. If a complex node has more than one subregion, the arity and order of all the internal inputs/outputs must match between all subregions, as well as match the arity and order of the external inputs/outputs of the complex node.

The complex nodes of an RVSDG relevant for this report are as follows:

- **γ -nodes: N-way statements**

γ -nodes represent conditional statements. Each γ -node has a predicate as first input. All other edges passing as inputs to the γ -node are edges its subregions depend upon. Each subregion represents one case. All subregions must have the same order and arity of internal inputs and outputs, even if the subgraph in each region does not depend on all of the internal outputs.

A γ -node is equivalent to a *switch-case* without fall-through in C/C++. Each case of the switch statement corresponds to a subregion of the γ -node. Hence, a simple *if-statement*

with no else-clause can be represented by a γ -node with two subregions. The true subregion contains the RVSDG subgraph that represents the body of the if-statement, whereas the false subregion of the γ -node simply routes all inputs through. See Figure 4 for an example of a γ -node.




Figure 4: Example of an RVSDG subgraph on the right, depicting a γ -node equivalent to the C/C++ if-statement on the left.

- **θ -nodes: Tail-controlled loops**

θ -nodes represent tail controlled loops. As with γ -nodes, its inputs (and outputs) are all the dependencies needed for the RVSDG subgraph in its subregion.

The first time the body of the loop is executed, the external inputs are mapped to the internal outputs, as Figure 5 exemplifies. This enables the complex node's contained RVSDG subgraph to execute with the values given as external inputs to the θ -node.

However, inside the θ -node there is an extra first internal input, which is the predicate of the tail controlled loop. If this predicate evaluates to true, the rest of the internal inputs of the θ -node are mapped to their corresponding internal outputs. This enables the iterative behaviour of an RVSDG θ -node. Thus, the operations represented by its contained RVSDG subgraph are executed as a tail-controlled loop. Finally, when the predicate evaluates to false, the internal inputs are mapped to the external outputs of the θ -node instead of the internal outputs.

A θ -node is equivalent to a *do-while* loop in C/C++, as shown in Figure 5.

```

unsigned int i = 0;
unsigned long r = 1;
do {
    i += 1;
    r *= i;
} while(i < n);

```




Figure 5: An RVSDG subgraph on the right, depicting a θ -node equivalent to the C/C++ do-while loop on the left.

A *head-controlled*-loop can be represented by putting a θ -node inside of the *true* clause of a γ -node containing no nodes in the subregion representing the *false* clause, as long as the γ - and θ -nodes share the same predicate.

- **λ -nodes: Functions**

λ -nodes represent functions. A λ -node contains an RVSDG subgraph representing the body of a function. The internal inputs of a λ -node represents the results of the function. Respectively, its internal outputs represent the arguments of the function. While λ -nodes don't have any external inputs, their external output are what give the *apply*-nodes their first input, enabling them to invoke the function represented by the λ -node.

The arity and order of a λ -node's internal inputs and outputs must match the arity and order of the external inputs and outputs of all connected *apply*-nodes.

Hence, when an *apply*-node connected with a λ -node is executed, the external inputs of the *apply*-node are mapped to the internal outputs of the λ -node, before the λ -node is invoked.

Likewise the internal inputs of the λ -node are gated through to the external outputs of the *apply*-node when the operations of the λ -node have been executed. See Figure 6 for an example of a λ -node.

```

unsigned long long
fac(unsigned int n){
    unsigned int i = 0;
    unsigned long r = 1;
    do{
        i += 1;
        r *= i;
    } while(i < n);
    return r;
}

```




Figure 6: Example of an RVSDG subgraph on the right, depicting a λ -node equivalent to the C/C++ function on the left.

- **ϕ -regions: Recursive environments**

ϕ -regions represent recursive environments. They contain at least one recursive λ -node. Like the λ -node, they have no external inputs. However, the internal outputs of the ϕ -region represent the links utilized by the *apply*-nodes contained within to connect with the respective λ -nodes.

The internal inputs of a ϕ -region receive the function invocation links from the λ -nodes contained within. The internal inputs map to the external outputs, thus enabling *apply*-nodes outside of the recursive environment to connect with the λ -nodes, as depicted in Figure 7.

```
unsigned long long
fac(unsigned int n){
    unsigned long r;
    if(n > 1){
        r = n*fac(n-1);
    } else{
        r = 1;
    }
    return r;
}
```




Figure 7: Example of an RVSDG subgraph on the right, depicting a ϕ -region equivalent to the recursive C/C++ function on the left.

3 The Inliner

The following heuristic is executed by the inliner of this project when given an RVSDG as input:

1. For all recursive environments (ϕ -regions):
 - (a) Use the approach described in Section 3.1 to fill a list *loop breakers*. These λ -nodes are *not* to be inlined.
2. Scan through the RVSDG, finding all *apply*-nodes. Exclude all function calls to loop breakers, calls invoking functions that are not-statically known, or external functions.
3. Make a list of the *apply*-nodes found in Step 2, and order the list of according to the heuristics discussed in Section 3.2. The order of *apply*-nodes inlined can affect the amount of *apply*-nodes inlined, even when each *apply*-node is evaluated with the same heuristic.
4. Look at each *apply*-node in turn from the list made in Step 3 and decide whether or not to inline it according to the heuristic discussed in Section 3.3:
 - (a) If the *apply*-node is inlined, add any newly copied (inlined) *apply*-nodes, following the same criteria as used in Step 2, to the list of *apply*-nodes. Continue with Step 3.
 - (b) If the *apply*-node is not inlined, continue with Step 4, evaluating the next *apply*-node.
5. When the inliner reaches the end of the list, no more *apply*-nodes have been inlined, and the inliner is finished.

3.1 Deciding which recursive functions to inline

The inliner visits all *apply*-nodes and the functions they invoke functions, with the same heuristic, described in Section 3.3 regardless of whether the functions invoked are recursive or not. However, the inliner of this project only considers inlining *some* of the *apply*-nodes invoking recursive functions, to ensure termination of the compiler.

The *apply*-nodes not considered for inlining are the ones invoking *loop breakers*. Loop breakers are function-nodes found in ϕ -regions which break the Strongly Connected Component cycles within the ϕ -region's call dependency graph.

Loop breakers are found by visiting each function in the recursive environment of a ϕ -region, and then checking what functions its *apply*-nodes call. If one of the *apply*-node calls a function already visited, the already visited function is marked as a loop breaker. Hence, all self-recursive functions are marked as loop breakers, but if two functions call each other in a cycle, only one of them may be added to the list of loop breakers.

Make a figure or three of ϕ -regions, one with a self-recursive lambda, one with some mutually recursive ones (the example from PVV w/Torje), and perhaps also one based on GHC papers FGHPQ example.

Hence, the inliner has a list of recursive functions which it knows *not* to inline, thus ensuring termination of the compilation. All other remaining recursive functions may then be safely inlined with the same criteria as any non-recursive functions.

3.2 The order of call sites inlined

The heuristic deciding on whether or not to inline a specific *apply*-node is based on previous work of Keith D. Cooper et. al [5] and Waterman [12]. Their work utilizes *Inlining Conditions* (ICs), which evaluate properties of the function invoked or the call-site(s) invoking function.

Because of above paragraph, which I think belongs to Section 3.3, ask Nico if he insists that Section 3.2 should come before Section 3.3. I feel that this makes little sense since the one that is first now relies on the second...

As such, when a successive series of functions call one another, our approach only considers one function's properties or *apply*-node(s)'s properties, depending upon the ICs used in the inlining heuristic, at a time. Some ICs such as *Node Count* (NC) can affect the amount of inlined *apply*-nodes, depending on the order the *apply*-nodes are visited in.

Given the criteria for whether an *apply*-node is inlined is that the function it invokes needs to contain less than four nodes, all of the *apply*-nodes in Figure 8 can be inlined, or just two of them. The order the *apply*-nodes are visited in is what makes the difference in the amount inlined given the inlining condition $NC < 4$.

When the inliner visits the *apply*-nodes in Figure 8 top down, λ_1 can be inlined into λ_2 . The new λ_{1-2} can be inlined into λ_3 , resulting in all of them being inlined into a single function λ_{1-2-3} . However, if the *apply*-nodes are visited by the inliner in a bottom up order instead, then λ_3 can be inlined into λ_2 , but λ_{2-3} 's NC exceeds 3, and cannot be inlined into λ_1 .

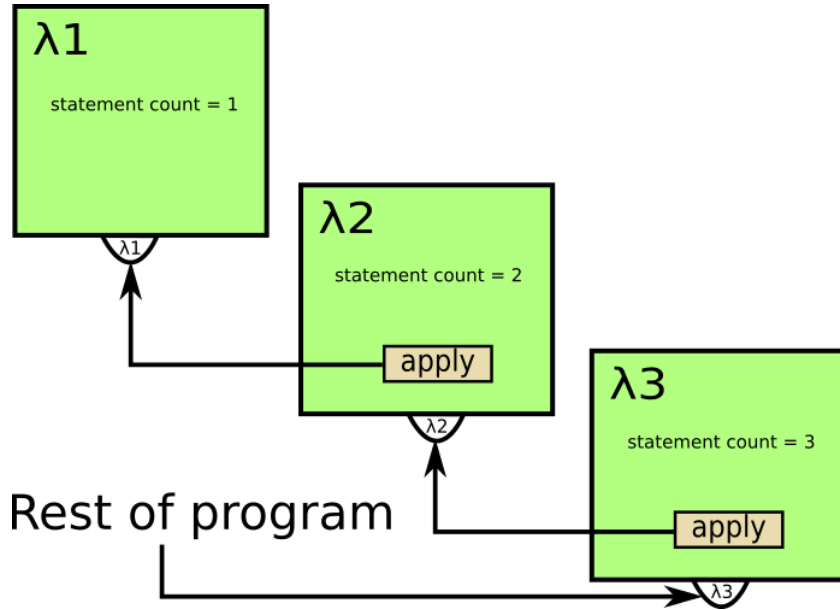


Figure 8: A minimal example of an RVSDG subgraph, depicting a function call order in a program.

Both the top down and bottom up examples inlining λ_1 , λ_2 , and λ_3 , suppose that no other optimizations are applied on the inlined code, meaning that no node which has gotten an *apply*-node inlined, will have less nodes than the sum of its original NC plus the NC of the function

inlined.

Because of this, our inliner is able to traverse the *apply*-nodes of an RVSDG in a top down order, or a bottom up order. Section 8.3 discusses ideas for potential work for future research on the impact of inlining call sites order.

3.3 Inlining a call site

Our approach was chosen because it permits effective testing for an apt heuristic when deciding on whether or not to inline a call site.

As mentioned, the inliner uses several different ICs: NC, *Static Call Count* (SCC), and *Loop Nesting Depth* (LND) being among these. SCC is the total amount of *apply*-nodes for the function invoked, and LND is how many nested loops the *apply*-node is located within.

These ICs and others described in Section 4.1, allow us to write and re-write the inlining heuristic effectively, by letting us write them using *Conjunctive Normal Form* (CNF). Thus, the CNFs written for our inliner heuristic are written in the following fashion: $NC < X \parallel SCC < Y \parallel (SCC < Z \ \&\& \ LND > W)$.

This permits us to efficiently search the parameter space for optimal parameters for the inlining heuristics.

Describe the algorithm and inliner conditions we land for evaluating a call site on after testing.

4 Methodology

Need an introduction here, no?

4.1 The Inlining Conditions (ICs)

The ICs utilized in this project are the following:

- **Node count (SN)**

This function property equates to the number of C/C++ statements contained within a function. A function's node count is an inliner condition we want to utilize because it gives us an idea of the size of the code-duplication if we inline the function.

- **Loop nesting depth (LND)**

This property tells us how potentially useful it is to inline this specific call site. The assumption is that most of a program's execution time is spent within loops, so there is potentially more to gain if optimizations are unveiled by inlining call sites inside nested loops.

- **Static call count (SCC)**

This property tells us how many *apply*-nodes invoke this function in the program. If this count is low, it may be worth inlining all the call sites and eliminating the original function. If the count is 1, then the call site can always be inlined, seeing as there is no risk of code-duplication.

- **Parameter count (PC)**

The greater the amount of parameters a function has, the greater the invocation cost of said function. This is especially true when type conversion is required. In some cases, the computational cost of an inlined with low node count may be smaller than the cost of invoking it if it has many parameters [12].

- **Constant parameter count (CPC)**

This property tells us how many of the call site's parameters are constant at the call site. Function invocations with constant parameters can often benefit more from unveiled optimizations after inlining.

- **Calls in procedure (CP)**

This function property tells us how many call sites are located inside the function the call site invokes. Hence, it enables finding leaf functions. Waterman [12] introduced this parameter for two distinct reasons: leaf functions are often small and easily inlined, and a high percentage of total execution time is spent in leaf functions.

5 Results

Make a subsection or paragraph showing the results of *apply*-nodes which link to static calls, vs all calls.

6 Related Work

As mentioned in Section 1, compilers have existed, and optimized code, since the last half of the 20th century. Inlining has long been an important optimization for most compilers. W. Davidson and M. Holler [6] examine the hypothesis that the increased code size of inlined code affects execution time on demand-paged virtual memory machines. Using equations developed to describe the execution time of an inlined program, they test this hypothesis through the use of a source-to-source subprogram inliner.

Cavazos and F.P. O’Boyle [3] use a genetic algorithm in their auto-tuning heuristics to show how conjunctive normal form (CNF) can easily be used to decide if and when to inline a specific call site. They report between 17% and 37% execution time improvements without code size explosion.

Serrano [11] implements an inliner in the Scheme programming language. The paper details an heuristic for which functions to inline, as well as an algorithm for how to inline recursive functions. The paper reports an average run time decrease of 15%.

Waterman’s Ph.D. thesis [12] examines the use of adaptive compilation techniques in combination with an inlining heuristic. His thesis shows how CNF can be used for deciding which functions to inline. It also details how there can be no single given correct set of parameters for all programs, given the search space of the heuristics hillclimbing algorithm. The thesis reports consistently better or equal run time compared to the GCC inliner and ATLAS.

D. Cooper et. al [5] expand on Waterman’s PhD Thesis [12]. Their paper details how the proper use of the parameterization search space using a hillclimber algorithm, in an adaptive inlining scheme, can achieve improved results compared to GCCs inliner. Their results range from 4% to 31% run time decrease compared to GCCs inliner.

E. Hank et. al [9] introduce a new technique called *Region-Based Compilation*. They examine the benefits an aggressive compiler gains from inlining on Very Long Instruction Word (VLIW) architectures. The paper reports that aggressive inlining can become costly, with an average code size expansion of 400%. However, their results also show that inlining is sufficiently able to unveil further compiler optimizations. Thus leading to an average of 50% of program execution time spent in functions with more than 1000 operations. This is an improvement, compared to their test results where more than 80% of the execution time was spent inside functions with less than 250 operations, when no inlining was employed.

P. Jones and Marlow [10] describe the inliner for the Glasgow Haskell Compiler (GHC). Their paper introduces a novel approach for deciding which mutually recursive functions can safely be inlined without code size explosion or the risk of non-termination. Jones and Marlow report on average of 30% run time decrease.

The report of Barton et. al [2] tests whether the potential for loop fusion should be taken into consideration in the inliner. They disprove this using the IBM®XL Compile Suite, measuring how many additional loops they were able to fuse in the SPECint2000 and SPECfp2000 benchmark suites. The results reported indicate that the compiler already catches most of the potential loop fusion optimizations, and the results cannot justify an inter-procedural loop fusion implementation.

Deshpande and A. Edwards [8] detail an inlining algorithm meant to improve inlining in the GHC. The algorithm improved the parallelism of recursive functions by “widening” them into the equivalent of multiple recursive calls through unrolling recursion. No results were reported.

W. Hwu and P. Chang [4] explore how program profile information could be used to decide

whether or not to statically inline C functions. Their motivation was to remove costly function calls in a C program, in addition to unveil potential optimizations. Through the use of the IMPACT-I C compiler, they profile dynamic program information, resulting in a call graph with weighted edges. They report 0% to 99% reduction of dynamic function calls in their test benchmarks.

7 Conclusion

8 Further Work

8.1 Dynamic profiling and adaptive compilation

Basically re-iterate Waterman's PhD idea.

8.2 Choosing loop breakers more carefully

Bas' MSc. idea

8.3 The ordering of which call sites to order first

Lots to say on ordering of call sites to be inlined

8.3.1 Something

If time permits, I want to briefly discuss the idea of evaluating all the apply nodes (in order), before inlining any. So as to be able to see the total cost of inlining a successive chain of function calls in one direction or the other.

List of Figures

1	Example of an RVSDG subgraph equivalent to the C/C++ on the left.	6
2	Example of an RVSDG subgraph equivalent to the C/C++ on the left.	7
3	A minimal example of a complex node, showing which inputs/outputs are external/internal, and how they can have multiple subregions.	8
4	Example of an RVSDG subgraph on the right, depicting a γ -node equivalent to the C/C++ if-statement on the left.	9
5	An RVSDG subgraph on the right, depicting a θ -node equivalent to the C/C++ do-while loop on the left.	10
6	Example of an RVSDG subgraph on the right, depicting a λ -node equivalent to the C/C++ function on the left.	11
7	Example of an RVSDG subgraph on the right, depicting a ϕ -region equivalent to the recursive C/C++ function on the left.	12
8	A minimal example of an RVSDG subgraph, depicting a function call order in a program.	14

Listings

1	C/C++ code showing the definitions of A() and B() when exemplifying CF. . . .	4
2	C/C++ code showing the definitions of C() and D(), when exemplifying code-duplication.	4

9 References

- [1] Helge Bahmann, Nico Reissmann, Magnus Jahre, and Jan Christian Meyer. Perfect reconstructability of control flow from demand dependence graphs. *ACM Trans. Archit. Code Optim.*, 11(4):66:1–66:25, January 2015.
- [2] Christopher Barton, José Nelson Amaral, and Bob Blainey. Should potential loop optimizations influence inlining decisions? In *Proceedings of the 2003 Conference of the Centre for Advanced Studies on Collaborative Research*, CASCON '03, pages 30–38. IBM Press, 2003.
- [3] John Cavazos and Michael F. P. O’Boyle. Automatic tuning of inlining heuristics. In *Proceedings of the 2005 ACM/IEEE Conference on Supercomputing*, SC '05, pages 14–, Washington, DC, USA, 2005. IEEE Computer Society.
- [4] P. P. Chang and W.-W. Hwu. Inline function expansion for compiling c programs. *SIGPLAN Not.*, 24(7):246–257, June 1989.
- [5] Keith D. Cooper, Timothy J. Harvey, and Todd Waterman. An adaptive strategy for inline substitution. In *Proceedings of the Joint European Conferences on Theory and Practice of Software 17th International Conference on Compiler Construction*, CC’08/ETAPS’08, pages 69–84, Berlin, Heidelberg, 2008. Springer-Verlag.
- [6] J.W. Davidson and A.M. Holler. Subprogram inlining: a study of its effects on program execution time. *Software Engineering, IEEE Transactions on*, 18(2):89–102, Feb 1992.
- [7] Bas den Heijer. Optimal loop breaker choice for inlining. Master’s thesis, Utrecht University, Netherlands, 2012.
- [8] Neil Ashish Deshpande and Stephen A Edwards. Statically unrolling recursion to improve opportunities for parallelism. 2012.
- [9] Richard E. Hank, Wen-Mei W. Hwu, and B. Ramakrishna Rau. Region-based compilation: An introduction and motivation. In *Proceedings of the 28th Annual International Symposium on Microarchitecture*, MICRO 28, pages 158–168, Los Alamitos, CA, USA, 1995. IEEE Computer Society Press.
- [10] Simon Peyton Jones and Simon Marlow. Secrets of the glasgow haskell compiler inliner. *J. Funct. Program.*, 12(5):393–434, July 2002.
- [11] Manuel Serrano. Inline expansion: When and how? In *Proceedings of the 9th International Symposium on Programming Languages: Implementations, Logics, and Programs: Including a Special Trach on Declarative Programming Languages in Education*, PLILP '97, pages 143–157, London, UK, UK, 1997. Springer-Verlag.
- [12] Todd Waterman. *Adaptive Compilation and Inlining*. PhD thesis, Houston, TX, USA, 2006. AAI3216796.

A Project Description

An Inliner for the Jive compiler

Nico Reissmann

Friday 12th December, 2014

Compilers have become an essential part of every modern computer system since their rise along with the emergence of machine-independent languages at the end of the 1950s. From the start, they not only had to translate between a high-level language and a specific architecture, but had to incorporate optimizations in order to improve code quality and be a par with human-produced assembly code. One such optimization performed by virtually every modern compiler is *inlining*. In principle, inlining is very simple: just replace a call to a function by an instance of its body. However, in practice careless inlining can easily result in extensive *work* and *code duplication*. An inliner must therefore decide carefully when and where to inline a function in order to achieve good performance without unnecessary code bloat.

The overall goal of this project is to implement and evaluate an inliner for the Jive compiler back-end. The project is split in a practical and an optional theoretical part. The practical part includes the following:

- Implementation of an inliner for the Jive compiler back-end. The inliner must be able to handle recursive functions and allow for the configuration of different heuristics to permit rapid exploration of the parameter space.
- An evaluation of the implemented inliner. A particular emphasis is given to different heuristics and their consequences for the resulting code in terms of work and code duplication.

The Jive compiler back-end uses a novel intermediate representation (IR) called the Regionalized Value State Dependence Graph (RVSDG). If time permits, the theoretical part of the project is going to clarify the consequences of using the RVSDG along with an inliner. It tries to answer the following research questions:

- What impact does the RVSDG have on the design of an inliner and the process of inlining?
- Does the RVSDG simplify/complicate the implementation of an inliner and the process of inlining compared to other commonly used IRs?

The outcome of this project is threefold:

1. A working implementation of an inliner in the Jive compiler back-end fulfilling the aforementioned criteria.
2. An evaluation of the implemented inliner.
3. A project report following the structure of a research paper.